CAISA™

Level 1 — AI Foundations Architect

CHAPTER 2 OF 48

Part I: ML Foundations · Week 1

Supervised Learning Concepts

Enterprise ML in Practice: Europe · UK · US · Canada

- **Study Time:** 2.0 Hours
- **Format:** Reading + Examples + Concept Checks
- **Jurisdictions:** EU · UK · US · Canada
- **Curriculum Position:** Week 1, Chapter 2 of 48

10 LEARNING OBJECTIVES INCLUDING:

- Define supervised learning and distinguish it from other ML paradigms
- Identify regression vs. classification and select the correct family
- Describe the 7-stage supervised ML pipeline end-to-end
- Explain bias–variance tradeoff and select evaluation metrics by business cost
- Analyse 5 enterprise case studies: ING, NHS, BNP Paribas, TD Bank, Siemens

CAISA™
Level 1 — AI Foundations Architect

CHAPTER 2 OF 48
Part I: ML Foundations · Week 1

Supervised Learning Concepts

Enterprise ML in Practice: Europe · UK · US · Canada

- **Study Time:** 2.0 Hours
- **Format:** Reading + Examples + Concept Checks
- **Jurisdictions:** EU · UK · US · Canada
- **Curriculum Position:** Week 1, Chapter 2 of 48

10 LEARNING OBJECTIVES INCLUDING:

- Define supervised learning and distinguish it from other ML paradigms
- Identify regression vs. classification and select the correct family
- Describe the 7-stage supervised ML pipeline end-to-end
- Explain bias–variance tradeoff and select evaluation metrics by business cost
- Analyse 5 enterprise case studies: ING, NHS, BNP Paribas, TD Bank, Siemens

Table of Contents

Learning Objectives

By the end of this chapter you will be able to:

1. Define supervised learning and distinguish it from unsupervised and reinforcement learning paradigms.

2. Identify the two fundamental problem families — regression and classification — and select the appropriate family for a given business scenario.

3. Describe the lifecycle of a supervised ML pipeline from raw data ingestion to production deployment.

4. Articulate the concept of a hypothesis function and explain how a model learns from labelled examples.

5. Explain the role of the loss function and how gradient-based optimisation minimises it.

6. Interpret training, validation, and test splits and justify their use in preventing data leakage.

7. Recognise at least eight canonical supervised learning algorithms and map each to its appropriate enterprise context.

8. Explain the bias–variance tradeoff and describe practical techniques for managing it.

9. Select evaluation metrics appropriate to the business cost structure of a given ML problem.

10. Analyse real-world supervised learning deployments across Europe (EU/UK), the US, and Canada.

■ Chapter Focus Statement

Supervised learning is the dominant ML paradigm in regulated enterprise environments. This chapter builds the conceptual scaffolding — labelled data, hypothesis functions, loss minimisation, generalisation — before the algorithm-specific chapters that follow. Every enterprise AI architect must be fluent in this vocabulary before engaging regulators, data scientists, or technology vendors.

2.1 What Is Supervised Learning?

Supervised learning is the branch of machine learning in which an algorithm is trained on a dataset of labelled examples — pairs (x, y) where x is a vector of input features and y is the known target output — with the goal of learning a generalised mapping function f such that f(x) ≈ y for previously unseen inputs.

The word 'supervised' is an analogy: the labelled training data acts as a teacher that provides correct answers, allowing the algorithm to adjust its internal parameters until its predictions are sufficiently accurate. Once trained, the model operates autonomously on new, unlabelled data — a capability at the heart of every production AI system in enterprise.

2.1.1 The Fundamental Learning Problem

Formally, supervised learning seeks to find a function f from a hypothesis space H that minimises the expected loss over the joint distribution P(X, Y):

Core Objective Function

f* = argmin(f ∈ H) E[L(y, f(x))] Where: L = loss function | f(x) = model prediction | y = true label E = expectation over the data distribution P(X, Y) In practice we minimise empirical risk — the average loss on the training set — ensuring that generalisation to unseen data is preserved through regularisation and cross-validation.

2.1.2 Supervised vs. Unsupervised vs. Reinforcement Learning

Paradigm	Training Signal	Core Mechanism	Enterprise Example
Supervised	Labelled (X, Y) pairs	Minimise prediction error against known labels	Credit scoring (ING), fraud detection (Barclays)
Unsupervised	Unlabelled X only	Discover hidden structure or compress representations	Customer segmentation (IKEA), anomaly detection (Siemens)
Semi-supervised	Small labelled + large unlabelled	Leverage unlabelled data to improve supervised predictions	Medical imaging (Philips, Siemens Healthineers)
Self-supervised	Auto-generated pseudo-labels from X	Pretraining representations (e.g., BERT, GPT)	Document intelligence (Deutsche Bank, HMRC)
Reinforcement	Reward signal from environment	Agent learns policy by maximising cumulative reward	Trading (Citadel Europe), robotics (ABB)

2.2 Regression vs. Classification — The Two Problem Families

Every supervised learning problem belongs to one of two families determined by the nature of the target variable y. This architectural decision shapes loss function, output layer, evaluation metric, and deployment design.

2.2.1 Regression — Predicting Continuous Outputs

Regression problems require the model to predict a real-valued number. Quality is typically measured by distance-based metrics such as MSE or MAE. Standard loss functions include MSE (penalises large errors more heavily, sensitive to outliers), MAE (robust to outliers, less differentiable at zero), and Huber loss (combines properties of both).

Regression — When to Use

Use when the target y is a continuous quantity: • Property valuations (Rightmove AI, Hometrack UK — predicting £ price) • Demand forecasting (Carrefour, Lidl — predicting daily units sold) • Energy load prediction (National Grid ESO — predicting MW demand) • LTV (Lifetime Value) estimation (Vodafone, Deutsche Telekom) • Salary benchmarking models (Workday EU data centres)

2.2.2 Classification — Predicting Discrete Labels

Classification problems require the model to assign each input to one of a finite set of categories. Binary classification is the simplest case; multi-class extends to three or more mutually exclusive categories; multi-label allows multiple simultaneous labels.

Classification — When to Use

Use when the target y is a discrete category: • Loan default prediction (ING, Lloyds Banking Group — binary: default / no default) • Transaction fraud detection (Mastercard, Stripe — binary: fraud / legitimate) • Disease diagnosis (NHS AI Lab, Roche — binary or multi-class) • Customer churn prediction (BT Group, Orange — binary: churn / retain) • Document routing (HMRC, BNP Paribas — multi-class across departments)

2.2.3 Recognising Problem Type from Business Requirements

Business Question	Problem Type	Target Variable y	Primary Metric
Will this customer repay their mortgage?	Binary Classification	Default (1) / No Default (0)	AUC-ROC, Recall
What will quarterly revenue be?	Regression	€ Revenue (continuous)	MAE, RMSE
Which product category will this order belong to?	Multi-class Classification	Category label (A, B … N)	F1 (Macro)
How long until this machine fails?	Regression (Survival)	Days to failure (continuous)	MAE, C-Index
Is this X-ray showing pneumonia?	Binary Classification	Positive (1) / Negative (0)	Recall (Sensitivity)
What is the risk tier for this SME loan?	Ordinal Classification	Low / Medium / High / Very High	Ordinal accuracy, AUC
How many units will we sell next week?	Regression (Time Series)	Units sold (continuous, t+7)	MAPE, RMSE

2.3 The Supervised Learning Pipeline

Enterprise supervised learning projects follow a structured 7-stage pipeline. Understanding this pipeline end-to-end is a prerequisite for the AI architect, who must reason about failure modes, regulatory exposure, and maintenance costs at each stage.

Stage 1 — Business & Problem Framing

Translate the business objective into a mathematically tractable prediction problem. Define the prediction target, decision boundary, business cost of each error type, applicable regulatory regime, and minimum acceptable performance threshold for production.

European Context: EU AI Act (high-risk systems), GDPR Article 22, EBA GL/2020/06, FCA SS1/23, OSFI E-23 (Canada), OCC SR 11-7 (US).

Stage 2 — Data Acquisition & Labelling

Supervised learning requires labelled data — the most expensive and time-consuming component. Sources in European enterprise: core banking systems (Temenos, Finastra), CRM (Salesforce), ERP (SAP, Oracle), open data portals (ONS UK, Eurostat, Statistics Canada), and EHDS/FIDA data-sharing frameworks.

Labelling strategies: expert annotation, programmatic labelling (Snorkel, Label Studio), and weak supervision. Label quality determines the performance ceiling.

Stage 3 — Feature Engineering

Select, transform, and construct features from raw data. Critical techniques: normalisation/standardisation, categorical encoding (one-hot, target encoding, embeddings), missing value handling, temporal feature extraction (lags, rolling means, seasonality), and feature selection (L1-regularised importance, mutual information).

Feature engineering remains the highest-leverage activity in applied ML for tabular enterprise data.

Stage 4 — Train / Validation / Test Split

Enforce the three-way split before any modelling begins. Apply temporal ordering for time-series data. Validate on a held-out validation set, and reserve the test set for final unbiased evaluation only.

Data leakage — contamination of training with future information — is the leading cause of production ML failure. Common causes: post-default features, full-dataset normalisation before splitting, temporal boundary violations.

Stage 5 — Model Training & Optimisation

Adjust model parameters to minimise the loss function. Parametric models use gradient descent variants (SGD, Adam). Ensemble methods use tree induction with hyperparameter search (grid, random, Bayesian via Optuna/Hyperopt).

Hyperparameter choices must be justifiable to model risk governance committees under EBA GL/2020/06 and OSFI E-23.

Stage 6 — Evaluation, Validation & Governance

Out-of-time testing, population stability analysis, fairness auditing (GDPR, EU AI Act, ECOA/US, Equality Act 2010/UK), and model risk review. High-risk sectors require independent second-line validation.

Fairness metrics must be computed across protected characteristics per applicable regulation in the deployment jurisdiction.

Stage 7 — Deployment & Monitoring

Production obligations: data drift monitoring (inputs shift over time), concept drift (relationship between inputs and outputs changes), performance degradation tracking, regulatory compliance. Tools: Evidently AI, Arize, WhyLabs.

Retraining schedules must be documented in the model governance framework.

2.4 Core Supervised Learning Algorithms — An Architect's Reference

The eight core supervised learning algorithm families encountered in production deployments across Europe, the US, and Canada. Subsequent chapters cover each algorithm in depth.

Algorithm	Type	Strengths	Limitations	Typical Enterprise Use
Linear Regression	Regression	Interpretable, fast, baseline standard	Assumes linearity; sensitive to outliers	Revenue forecasting (Siemens, Unilever)
Logistic Regression	Classification	Probability outputs; audit-friendly	Cannot capture non-linear relationships	EU credit scoring (ING, Société Générale)
Decision Tree	Both	Highly interpretable; no scaling needed	Prone to overfitting; unstable	Regulatory rule extraction (BaFin, FCA)
Random Forest	Both	Robust; handles missing data well	Slower; less interpretable than trees	Fraud detection (Mastercard Europe)
Gradient Boosting (XGBoost)	Both	State-of-the-art on tabular data	Hyperparameter sensitive; longer training	Risk scoring (HSBC, BNP Paribas)
Support Vector Machine	Classification	Effective in high dimensions	Slow on large datasets; no probabilities	Document classification (law, insurance)
k-Nearest Neighbour	Both	Simple; no training phase	Very slow at inference; memory intensive	Anomaly detection (legacy fraud systems)
Naïve Bayes	Classification	Extremely fast; works with small data	Assumes feature independence	Email spam, NLP classification

Algorithm Selection Principle

Algorithm choice is rarely the primary driver of enterprise AI success. The dominant factors are: data quality and quantity, feature engineering, robust validation methodology, regulatory fit, and operational maintainability. An interpretable logistic regression that survives a regulator's audit is worth more to the business than a black-box ensemble that cannot be explained.

2.5 The Bias–Variance Tradeoff

Bias–Variance Decomposition

Expected Error = $Bias^2$ + Variance + Irreducible Noise $Bias^2$ = Error from incorrect model assumptions (underfitting) Variance = Error from sensitivity to training set fluctuations (overfitting) Irreducible = Error inherent in the data-generating process (cannot be reduced) The architect's goal is to find the regularised minimum of $Bias^2$ + Variance.

2.5.1 High Bias — Underfitting

A high-bias model has made overly simplistic assumptions about the data-generating process. It performs poorly on both training and test sets. Symptoms: training loss is high; training and validation loss are similarly high; residual plots show systematic patterns.

Causes: model too simple for the relationship, insufficient features, excessive regularisation, too few training iterations. Remediation: add relevant features, use a more expressive model class, reduce regularisation strength.

2.5.2 High Variance — Overfitting

A high-variance model has memorised the training data including its noise. It performs well on the training set but poorly on validation and test. Symptoms: very low training loss, high validation loss, large generalisation gap, predictions that fluctuate across different training subsets.

Remediation: more training data, L1/L2 regularisation, dropout (neural networks), cross-validation, reduced model complexity, early stopping, bagging ensembles (Random Forest).

2.5.3 Practical Bias–Variance Management

Condition	Training Error	Test Error	Architectural Response
High Bias (Underfit)	High	High	Add features, increase model complexity, reduce L2 penalty
High Variance (Overfit)	Low	High	Add L1/L2 regularisation, more data, dropout, pruning
Optimal	Low-Moderate	Low-Moderate	Cross-validate; monitor on held-out production slice
Data Leakage	Very Low	Very High (spike)	Audit feature pipeline; enforce temporal splits strictly

Cross-Validation in Regulated Environments

K-fold cross-validation (k=5 or k=10) is the standard method for estimating generalisation performance. In time-series settings — the norm in financial services, insurance, and retail across Europe and North America — standard k-fold is invalid because it allows future data to inform predictions about the past. Walk-forward (expanding window) or time-series split cross-validation must be used instead. This is a mandatory validation requirement under EBA GL/2020/06 and is examined in CAISA Level 2.

2.6 Evaluation Metrics — Selecting the Right Measure

Choosing the correct evaluation metric is one of the most consequential architectural decisions in supervised ML. The wrong metric can optimise for a quantity that does not reflect the actual business objective, leading to models that score well in development but deliver poor business outcomes in production.

Metric	Formula	When to Prioritise	Enterprise Example
Accuracy	Correct / Total	Balanced classes only	Basic benchmarking (rarely used alone)
Precision	TP / (TP + FP)	Cost of false positives is high	Anti-money laundering (HSBC, Santander)
Recall (Sensitivity)	TP / (TP + FN)	Cost of false negatives is high	Cancer screening (NHS AI, Philips Health)
F1 Score	2 × P × R / (P + R)	Imbalanced classes; balanced tradeoff	Fraud detection (Mastercard, Stripe)
AUC-ROC	Area under ROC curve	Model ranking & threshold selection	Credit risk (ING, TD Bank, Wells Fargo)
Log Loss	–log P(y|x)	Probability calibration essential	Insurance pricing (Zurich, AXA, Intact)
MCC	Balanced metric for imbalance	Extreme class imbalance	Rare disease prediction (Roche, Novartis)

2.6.1 Class Imbalance and Its Consequences

Enterprise ML datasets are frequently highly imbalanced. Fraud rates in European card payment networks are approximately 0.04–0.08% of transactions (ECB Payment Statistics, 2023). Insurance claim fraud detection datasets typically show 1–3% positive rates. Clinical trial datasets for rare diseases may have positive rates below 0.1%.

In these settings, accuracy is a misleading metric: a model that predicts 'no fraud' for every transaction achieves 99.96% accuracy while catching zero fraudulent transactions. Architects must use Precision, Recall, F1, AUC-ROC, and MCC instead, and must apply imbalance-handling strategies: oversampling (SMOTE), undersampling, class-weighted loss functions, or threshold calibration.

Threshold Calibration

> Most classifiers output a probability $P(y=1|x)$, not a binary decision. The decision threshold — the probability cutoff above which the model predicts the positive class — is a business parameter, not a model parameter. Lowering the threshold increases recall (catches more fraud) at the cost of precision (more false positives). The optimal threshold is determined by the relative business costs of each error type — a decision that must involve domain experts, not just data scientists.

2.7 Enterprise Case Studies

Five real-world supervised learning deployments in regulated enterprise environments. Primary focus is Europe, with complementary cases from the US and Canada. Each case analyses problem framing, algorithm choice, evaluation strategy, and regulatory considerations.

CASE STUDY

ING Group

■■ Netherlands — Europe

Retail Banking | Credit Risk

Context	ING Group, headquartered in Amsterdam, is one of Europe's largest retail and commercial banks with over 38 million customers across 40 countries, operating under Dutch National Bank (DNB) supervision and subject to ECB SREP requirements and the EU AI Act.
Business Problem	Automate mortgage affordability assessment for Dutch and Belgian retail customers, reducing manual underwriting time from 5–7 days to same-day decisions while maintaining regulatory compliance with DNB model risk guidelines.
ML Approach	Binary classification (default / no default, 24-month horizon) using Gradient Boosting (LightGBM). Features: income-to-debt ratio, employment type, LTV ratio, credit bureau scores, open banking transaction data (PSD2-enabled), and macroeconomic features (ECB rates, Dutch house price index, unemployment).
Regulatory Fit	Classified as high-risk under EU AI Act Annex III. Requires conformity assessment, Annex IV technical documentation, and human oversight. A loan officer review is triggered for borderline scores within ±0.08 of the decision boundary. Validated by second-line MRM function per ECB model risk expectations (2022).
Outcome	Decision time reduced from 5.7 days to 4.2 hours. Gini coefficient improved by 8 points vs. prior scorecard. 34% reduction in manually reviewed applications. Deployed via Azure ML with automated daily drift monitoring.

CASE STUDY

NHS AI Lab / Moorfields Eye Hospital

■■ United Kingdom — Europe

Healthcare | Medical Imaging

Context	The NHS AI Lab, in partnership with Moorfields Eye Hospital and Google DeepMind (now Google Health), deployed a supervised learning system for automated detection and referral of over 50 sight-threatening eye diseases from optical coherence tomography (OCT) scans.

Business Problem	Reduce the diagnostic backlog for specialist retinal referrals in NHS ophthalmology clinics, where average wait times exceed 66 days, threatening permanent sight loss in conditions such as age-related macular degeneration, diabetic macular oedema, and glaucoma.
ML Approach	Multi-class classification across 50+ disease presentations. Architecture: 3D U-Net for volumetric OCT segmentation followed by a classification head. Primary metric: Sensitivity (Recall) — because a missed diagnosis carries catastrophically higher cost than an unnecessary referral.
Regulatory Fit	Regulated as a Medical Device (Class IIb) under UK MDR 2002 and UKCA marking. Required MHRA approval with pre-market clinical evidence. All AI recommendations reviewed by a trained ophthalmic technician or nurse before action. Published in Nature Medicine (De Fauw et al., 2018).
Outcome	System matched the diagnostic performance of 8 leading ophthalmologists on the test set. Estimated NHS efficiency saving: freeing 36% of specialist consultation slots. Referral accuracy: 94.0% for most urgent tiers.

CASE STUDY

BNP Paribas

■■ **France — Europe**

Investment Banking | Trade Finance | AML

Context	BNP Paribas, Europe's largest bank by assets (€2.6 trillion AUM), deployed a supervised learning system within its Trade Finance division to automate screening of documentary credits for sanctions compliance and AML indicators.
Business Problem	Manual compliance review of 80,000+ daily trade finance transactions took 2.3 days per document set. False positives from rule-based systems were flagging 22% of transactions for manual review, of which fewer than 0.3% represented genuine risk.
ML Approach	Binary classification (suspicious / legitimate) using an ensemble of Logistic Regression (regulatory audit trail) and XGBoost (predictive power). NLP applied via TF-IDF on free-text LC document fields. Optimised for Recall ≥ 99.5% for high-risk corridors, with precision maximised within that constraint.
Regulatory Fit	Governed under BNP Paribas internal Model Risk Policy aligned with BCBS 239. Supervised by French ACPR. External audit by a Big Four firm annually. Explainability provided via SHAP values integrated into the compliance officer's decision screen.
Outcome	False positive rate reduced from 22% to 4.7% (79% reduction). Average review time reduced from 2.3 days to 6.2 hours. Estimated annual compliance cost saving: €28 million across European operations.

CASE STUDY

TD Bank Group

■■ Canada + ■■ United States

Retail Banking | Customer Lifetime Value

Context	TD Bank Group, headquartered in Toronto and operating extensively across Canada and the eastern US, deployed a supervised regression model to estimate Customer Lifetime Value (CLV) for personalised product recommendation and proactive relationship management across 27 million active customers.
Business Problem	Retail and commercial banking divisions lacked a unified, forward-looking customer value metric, leading to suboptimal resource allocation in relationship management, product cross-sell, and attrition intervention programmes.
ML Approach	Regression: predict total discounted margin contribution per customer over a 36-month horizon (continuous output in CAD$). XGBoost ensemble with separate models per customer lifecycle stage. Primary metric: MAE in CAD$ — chosen because the business cost of a prediction error is approximately linear with error magnitude.
Regulatory Fit	Governed under OSFI Guideline E-23 (Canada) and OCC SR 11-7 (US). Credit pricing use cases subject to US Equal Credit Opportunity Act (ECOA) fair lending review. Canadian OSFI examination includes explainability and stability assessments.
Outcome	26% improvement in targeted product offer acceptance rates. 19% reduction in churn among top-quintile CLV customers where proactive intervention was applied. Estimated incremental revenue: CAD $340 million over the first 18 months of deployment.

CASE STUDY

Siemens AG — Digital Industries

■■ Germany — Europe

Manufacturing | Predictive Maintenance

Context	Siemens AG, headquartered in Munich, deployed a supervised learning system within its Digital Industries division to predict motor and CNC machine failures at customer facilities, delivered as a SaaS offering via the Siemens MindSphere industrial IoT platform.
Business Problem	Unplanned machine downtime in European automotive and semiconductor manufacturing costs an estimated €864 billion annually (Senseye/McKinsey, 2022). Siemens sought to predict failures 72+ hours in advance as a commercial value-added service.
ML Approach	Binary classification (failure within 72 hours / no failure) using Random Forest with survival analysis features. FFT frequency bands, temperature sensors, power consumption, and motor current signature analysis. Primary metric: Recall ≥ 85% at ≤ 15% false positive rate.

Regulatory Fit	Subject to EU Machinery Regulation (2023/1230) for safety-critical decisions. Requires risk assessment, CE marking documentation. The system is advisory — human maintenance engineers make final intervention decisions.
Outcome	Average unplanned downtime reduced by 36% across deployed customer base. False alarm rate reduced from 38% to 11%. Customer retention in MindSphere Predictive Services tier: 91% annual renewal rate.

2.8 Chapter Summary

Supervised learning is the paradigm that underpins the vast majority of enterprise AI deployments in regulated industries across Europe, North America, and globally. Its power lies in its clarity: given labelled examples of past behaviour, the model learns to generalise, enabling organisations to make consistent, scalable predictions about future events — from mortgage defaults to retinal disease to industrial equipment failure.

The five case studies — spanning ING Group (Netherlands), the NHS AI Lab and Moorfields (UK), BNP Paribas (France), TD Bank (Canada / US), and Siemens Digital Industries (Germany) — demonstrate a consistent pattern: enterprise AI success is not primarily determined by algorithm sophistication, but by disciplined problem framing, regulatory-aware design, rigorous data governance, and continuous production monitoring.

Key Takeaways for the AI Architect

1. Supervised learning requires labelled data, a defined prediction target, and a measurable business outcome. 2. Regression vs. classification is the primary architectural decision — algorithm and metric selection follows. 3. The 7-stage pipeline must be followed rigorously. Shortcutting any stage introduces compounding risk. 4. The bias–variance tradeoff governs generalisation. Regularisation and cross-validation are the primary controls. 5. Evaluation metrics must reflect business cost structure. Accuracy is rarely the right metric with imbalanced classes. 6. In Europe: GDPR Art. 22, EU AI Act, EBA GL/2020/06, FCA SS1/23 apply. In Canada: OSFI E-23. In the US: OCC SR 11-7 and ECOA.

Further Reading & References

• Hastie, T., Tibshirani, R., & Friedman, J. (2009). The Elements of Statistical Learning. Springer. [Chapters 2, 3, 4, 7, 10]

• James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021). An Introduction to Statistical Learning (2nd ed.). Springer. [Free PDF: statlearning.com]

• Chen, T. & Guestrin, C. (2016). XGBoost: A Scalable Tree Boosting System. Proceedings of KDD 2016.

• De Fauw, J. et al. (2018). Clinically applicable deep learning for diagnosis and referral in retinal disease. Nature Medicine, 24, 1342–1350.

• European Banking Authority (2020). EBA/GL/2020/06 — Guidelines on Loan Origination and Monitoring. EBA.

• European Central Bank (2022). Guide on Internal Models — General Topics. ECB Banking Supervision.

• Financial Conduct Authority / PRA (2023). SS1/23 — Model Risk Management Principles for Banks. Bank of England.

• Office of the Superintendent of Financial Institutions (2023). Guideline E-23: Model Risk Management. OSFI Canada.

• Office of the Comptroller of the Currency / Federal Reserve (2011). SR 11-7 — Supervisory Guidance on Model Risk Management.

• European Parliament (2024). Regulation (EU) 2024/1689 — Artificial Intelligence Act. Official Journal of the EU.

- Senseye / McKinsey (2022). True Cost of Downtime — Industrial Machinery in Europe. Senseye Research.
- Lundberg, S. & Lee, S.I. (2017). A Unified Approach to Interpreting Model Predictions (SHAP). NeurIPS 2017.

CAISA™ Level 1 — Chapter 2 of 48 | Part I: Machine Learning Foundations